

```
In[]: ##Name: Tanmay Gujar
      ##Rollno: T4159
      ##Subject: LP-IV(DL)

In[1]: import numpy as np
import pandas as pd
import random
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Conv2D, Dense, MaxPooling2D
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.datasets import mnist

In[2]: (X_train, y_train), (X_test, y_test) = mnist.load_data()

In[3]: print(X_train.shape)

(60000, 28, 28)

In[4]: X_train[0].min(), X_train[0].max()

Out[4]: (0, 255)

In[5]: X_train=(X_train-0.0)/(255.0-0.0)
X_test=(X_test-0.0)/(255.0-0.0)
X_train[0].min(), X_train[0].max()

Out[5]: (0.0, 1.0)

In[6]: def plot_digit(image, digit, plt, i):
plt.subplot(4, 5, i+1)
plt.imshow(image, cmap=plt.get_cmap('gray'))
plt.title(f'Digit: {digit}')
plt.xticks([])
plt.yticks([])
plt.figure(figsize=(16, 10))
for i in range(20):
plot_digit(X_train[i], y_train[i], plt, i)
plt.show()

Digit: 5
Digit: 0
Digit: 4
Digit: 1
Digit: 9
Digit: 2
Digit: 1
Digit: 3
Digit: 1
Digit: 4
Digit: 3
Digit: 5
Digit: 3
Digit: 6
Digit: 1
Digit: 7
Digit: 2
Digit: 8
Digit: 6
Digit: 9

In[7]: X_train=X_train.reshape((X_train.shape+(1,)))
X_test=X_test.reshape((X_test.shape+(1,)))

In[8]: y_train[0:20]

Out[8]: array([5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3, 5, 3, 6, 1, 7, 2, 8, 6, 9],
dtype=uint8)

In[9]: model=Sequential([
Conv2D(32, (3, 3), activation="relu", input_shape=(28, 28, 1)),
MaxPooling2D((2, 2)), Flatten(),
Dense(100, activation="relu"), Dense(10, activation="softmax")
])

In[10]: optimizer=SGD(learning_rate=0.01, momentum=0.9)
model.compile(optimizer=optimizer, loss="sparse_categorical_crossentropy", metrics=["accuracy"])
model.summary()

Model: "sequential"
Layer (type) OutputShape Param#
conv2d (Conv2D) (None, 26, 26, 32) 320
max_pooling2d (MaxPooling2D) (None, 13, 13, 32) 0
flatten (Flatten) (None, 5408) 0
dense (Dense) (None, 100) 540900
dense_1 (Dense) (None, 10) 1010
Total params: 542,230
Trainable params: 542,230
Non-trainable params: 0

In[11]: model.fit(X_train, y_train, epochs=10, batch_size=32)

Epoch 1/10
1875/1875 [=====] -19s 10ms/step - loss: 0.2392 - accuracy: 0.9275
Epoch 2/10
1875/1875 [=====] -18s 9ms/step - loss: 0.0801 - accuracy: 0.9757
Epoch 3/10
1875/1875 [=====] -18s 10ms/step - loss: 0.0511 - accuracy: 0.9846
Epoch 4/10
1875/1875 [=====] -18s 10ms/step - loss: 0.0352 - accuracy: 0.9893
Epoch 5/10
1875/1875 [=====] -18s 10ms/step - loss: 0.0268 - accuracy: 0.9917
Epoch 6/10
1875/1875 [=====] -18s 10ms/step - loss: 0.0203 - accuracy: 0.9938
Epoch 7/10
1875/1875 [=====] -18s 10ms/step - loss: 0.0145 - accuracy: 0.9955
Epoch 8/10
1875/1875 [=====] -18s 10ms/step - loss: 0.0107 - accuracy: 0.9969
Epoch 9/10
1875/1875 [=====] -19s 10ms/step - loss: 0.0092 - accuracy: 0.9975
Epoch 10/10
1875/1875 [=====] -18s 10ms/step - loss: 0.0066 - accuracy: 0.9983

Out[11]: <keras.callbacks.History at 0x27b921bdfa0>

In[12]: plt.figure(figsize=(16, 10))
for i in range(20):
image = random.choice(X_test).squeeze()
digit = np.argmax(model.predict(image.reshape((1, 28, 28, 1))))[0]
axis = 1
plot_digit(image, digit, plt, i)
plt.show()

1/1 [=====] -0s 105ms/step
1/1 [=====] -0s 19ms/step
1/1 [=====] -0s 18ms/step
1/1 [=====] -0s 13ms/step
1/1 [=====] -0s 20ms/step
1/1 [=====] -0s 20ms/step
1/1 [=====] -0s 20ms/step
1/1 [=====] -0s 18ms/step
1/1 [=====] -0s 19ms/step
1/1 [=====] -0s 19ms/step
1/1 [=====] -0s 18ms/step
1/1 [=====] -0s 19ms/step
1/1 [=====] -0s 19ms/step
1/1 [=====] -0s 20ms/step
1/1 [=====] -0s 20ms/step
1/1 [=====] -0s 20ms/step
1/1 [=====] -0s 20ms/step
1/1 [=====] -0s 21ms/step
1/1 [=====] -0s 20ms/step

Digit: 8
Digit: 9
Digit: 4
Digit: 8
Digit: 7
Digit: 2
Digit: 0
Digit: 3
Digit: 8
Digit: 8
Digit: 7
Digit: 3
Digit: 7
Digit: 9
Digit: 6
Digit: 7
Digit: 3
Digit: 2
Digit: 3
Digit: 0

In[13]: predictions = np.argmax(model.predict(X_test), axis=-1)
accuracy_score(y_test, predictions)

313/313 [=====] -1s 3ms/step
Out[13]: 0.9881

In[14]: n = random.randint(0, 9999)
plt.imshow(X_test[n])
plt.show()

0
5
10
15
20
25
0 5 10 15 20 25

In[15]: predicted_value = model.predict(X_test)
print(f"Handwritten number in the image is: {np.argmax(predicted_value[n])}")

313/313 [=====] -1s 3ms/step
Handwritten number in the image is: 0

In[16]: score = model.evaluate(X_test, y_test, verbose=0)
print(f'Test loss: {score[0]}')
print(f'Test accuracy: {score[1]}')

Test loss: 0.03916610777378082
Test accuracy: 0.98809992275238

In[17]: # The implemented CNN model is giving loss = 0.04624301567673683 and accuracy = 0.9872000217437744 for the test mnist dataset
```